

# EEP 596: LLMs: From Transformers to GPT || Lecture 4

Dr. Karthik Mohan

Univ. of Washington, Seattle

January 22, 2025

# Outline for Lecture

- Training and Back-propagation
- 

# Outline for Lecture

- Training and Back-propagation
  - Over-fitting and Hyper-parameters
- 

# Outline for Lecture

- Training and Back-propagation
  - Over-fitting and Hyper-parameters
  - Other DL architectures
- 

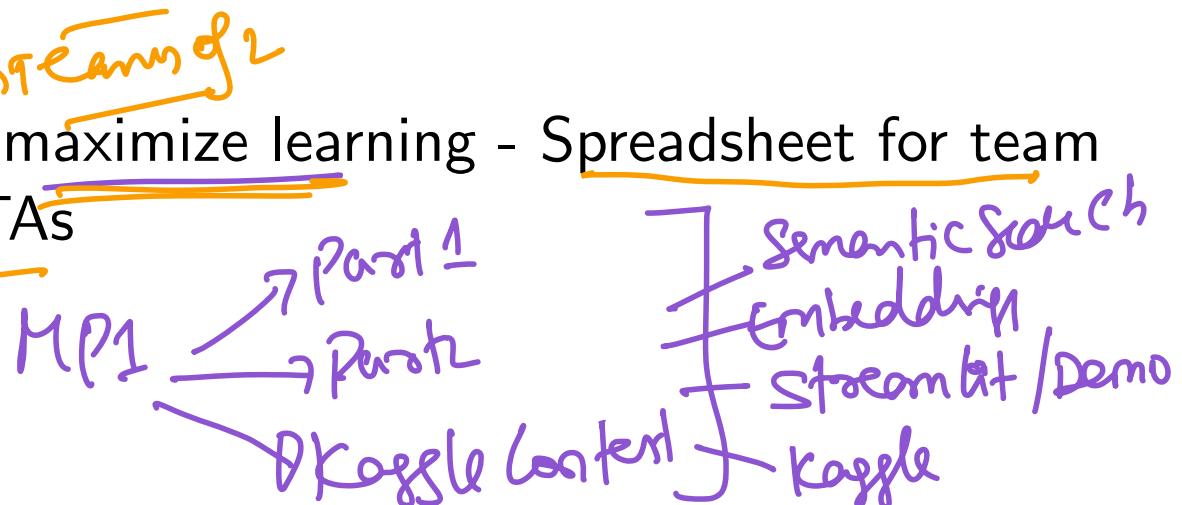
# Outline for Lecture

- Training and Back-propagation
- Over-fitting and Hyper-parameters
- Other DL architectures
- **Embeddings and Semantic Search**



# House Keeping Items

- Office Hours and Review Hours → webpage (GeoChip support)
- Assignment 2 will be a Mini-Project 1 assigned today → 2.5 weeks
- Mini-Project 1 will come in two main parts and cover embeddings and semantic search
- Two people per team to maximize learning - Spreadsheet for team pairing to be shared by TAs
- Any questions?



# Deep Learning Reference

## Deep Learning

Great reference for the theory and fundamentals of deep learning: Book by Goodfellow and Bengio et al [Bengio et al](#)

## Deep Learning History

# Recap from Last time!

- Perceptron and Logistic Regression

# Recap from Last time!

- Perceptron and Logistic Regression
  - How can 2-layer NN learn an XOR function
- 

# Recap from Last time!

- Perceptron and Logistic Regression
- How can 2-layer NN learn an XOR function
- Concepts of right fit, over-fitting and under-fitting

# Recap from Last time!

- Perceptron and Logistic Regression
- How can 2-layer NN learn an XOR function
- Concepts of right fit, over-fitting and under-fitting
- Tensorflow demo

# Multi-Layer Perceptron (MLP)

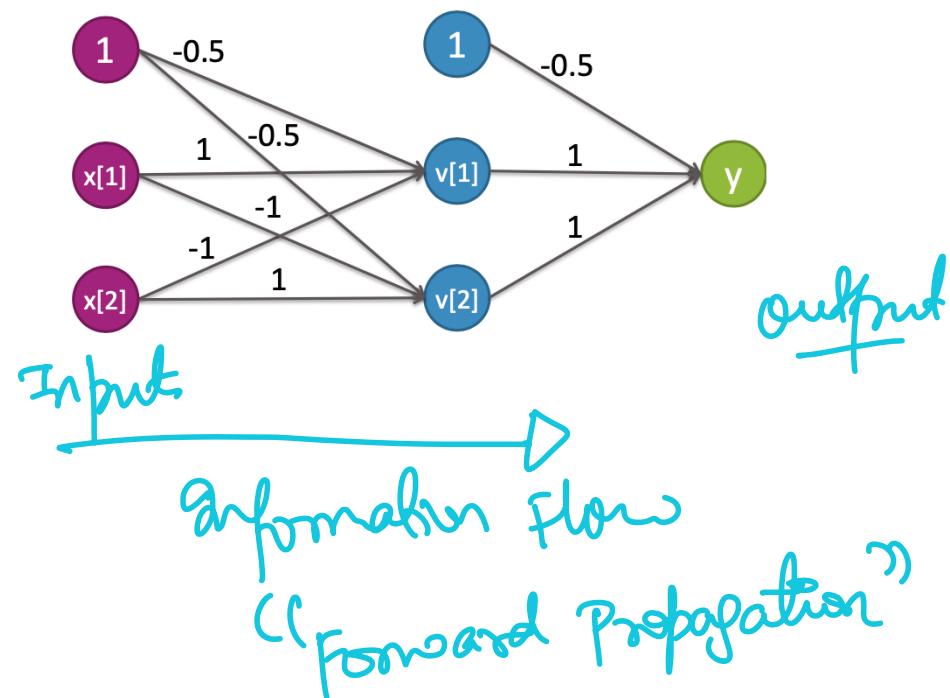
This is a 2-layer neural network

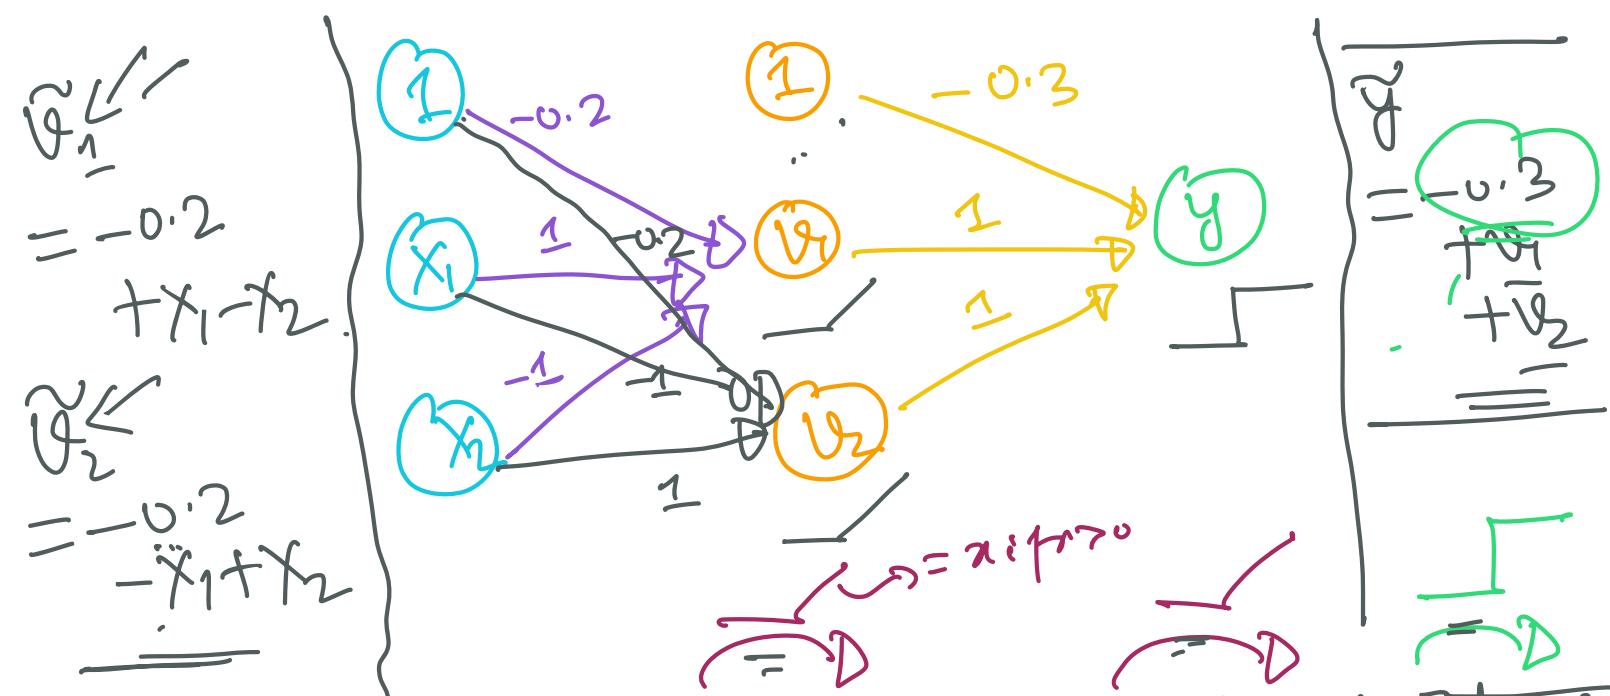
$$y = x[1] \text{ XOR } x[2] = (x[1] \text{ AND } !x[2]) \text{ OR } (!x[1] \text{ AND } x[2])$$

$$\begin{aligned}v[1] &= (x[1] \text{ AND } !x[2]) \\&= g(-0.5 + x[1] - x[2])\end{aligned}$$

$$\begin{aligned}v[2] &= (!x[1] \text{ AND } x[2]) \\&= g(-0.5 - x[1] + x[2])\end{aligned}$$

$$\begin{aligned}y &= v[1] \text{ OR } v[2] \\&= g(-0.5 + v[1] + v[2])\end{aligned}$$





$x_1$	$x_2$	$v_1$	$v_1$	$v_2$	$y$	$y$
0	0	-0.2	0	-0.2	0	-0.3
1	0	0.8	0.8	-1.2	0	0.5
0	1	-1.2	0	0.8	0.8	0.5
1	1	-0.2	0	-0.2	0	0.3

Below the table:

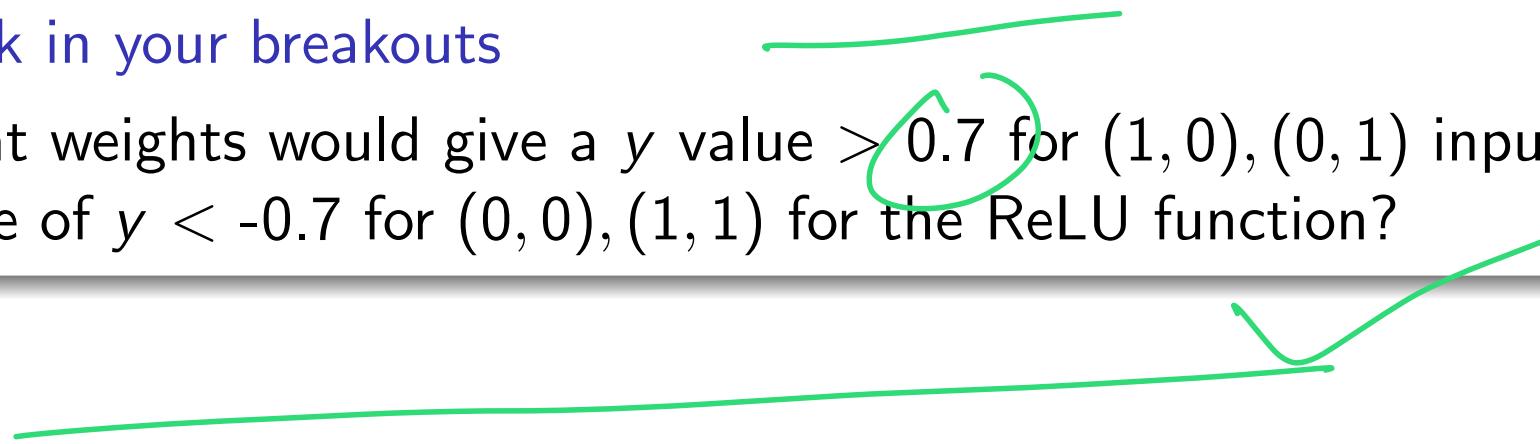
- $-0.2 + 1 - 1 = -0.2$
- $-0.2 - 1 + (1) = -0.2$

ICE from last lecture ✓

# Breakout Session: Would ReLU work?

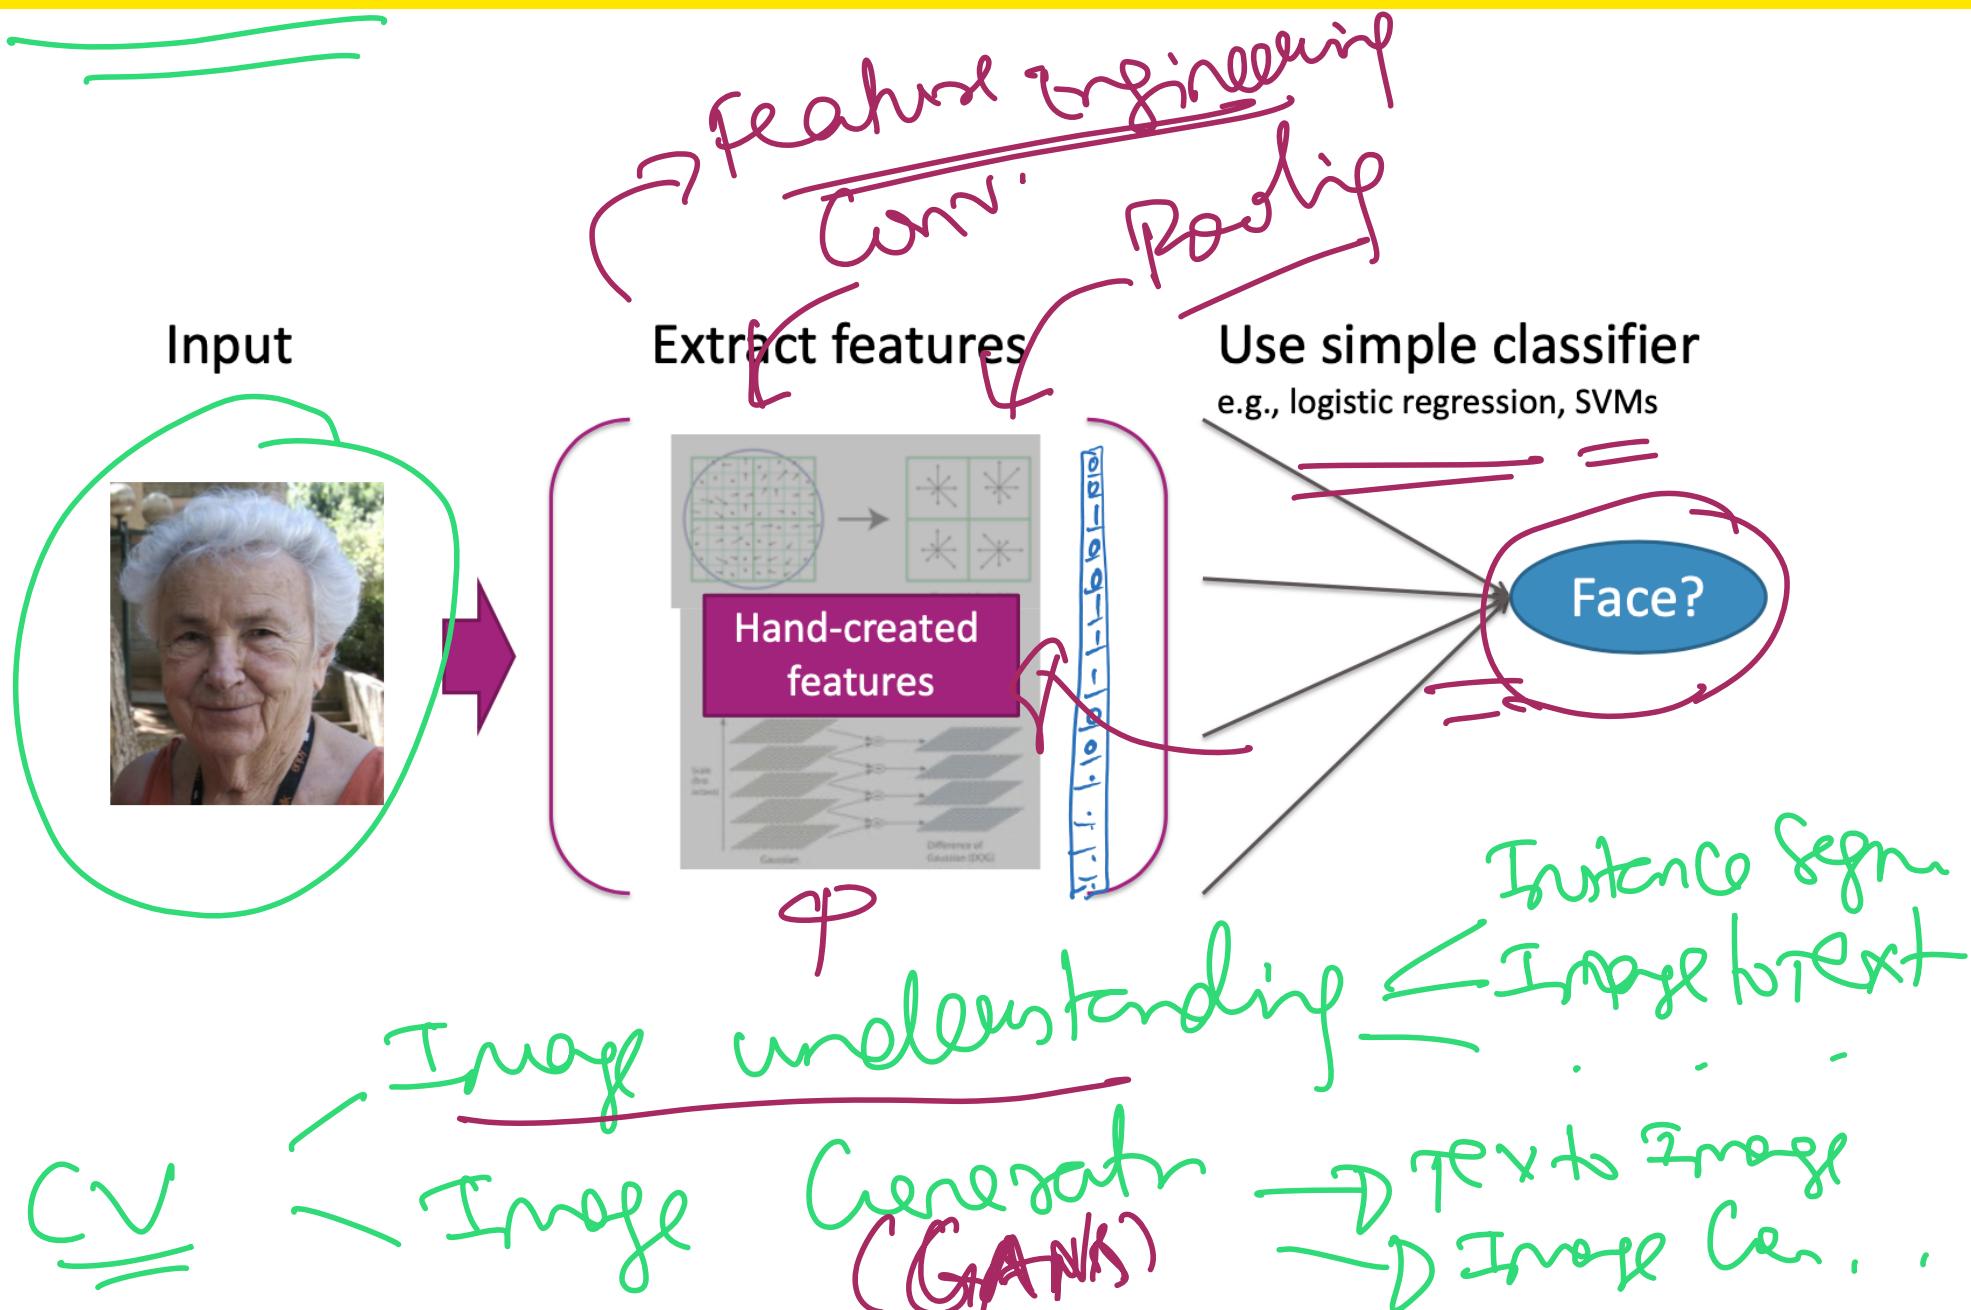
Work in your breakouts

What weights would give a  $y$  value  $> 0.7$  for  $(1, 0), (0, 1)$  inputs and a value of  $y < -0.7$  for  $(0, 0), (1, 1)$  for the ReLU function?

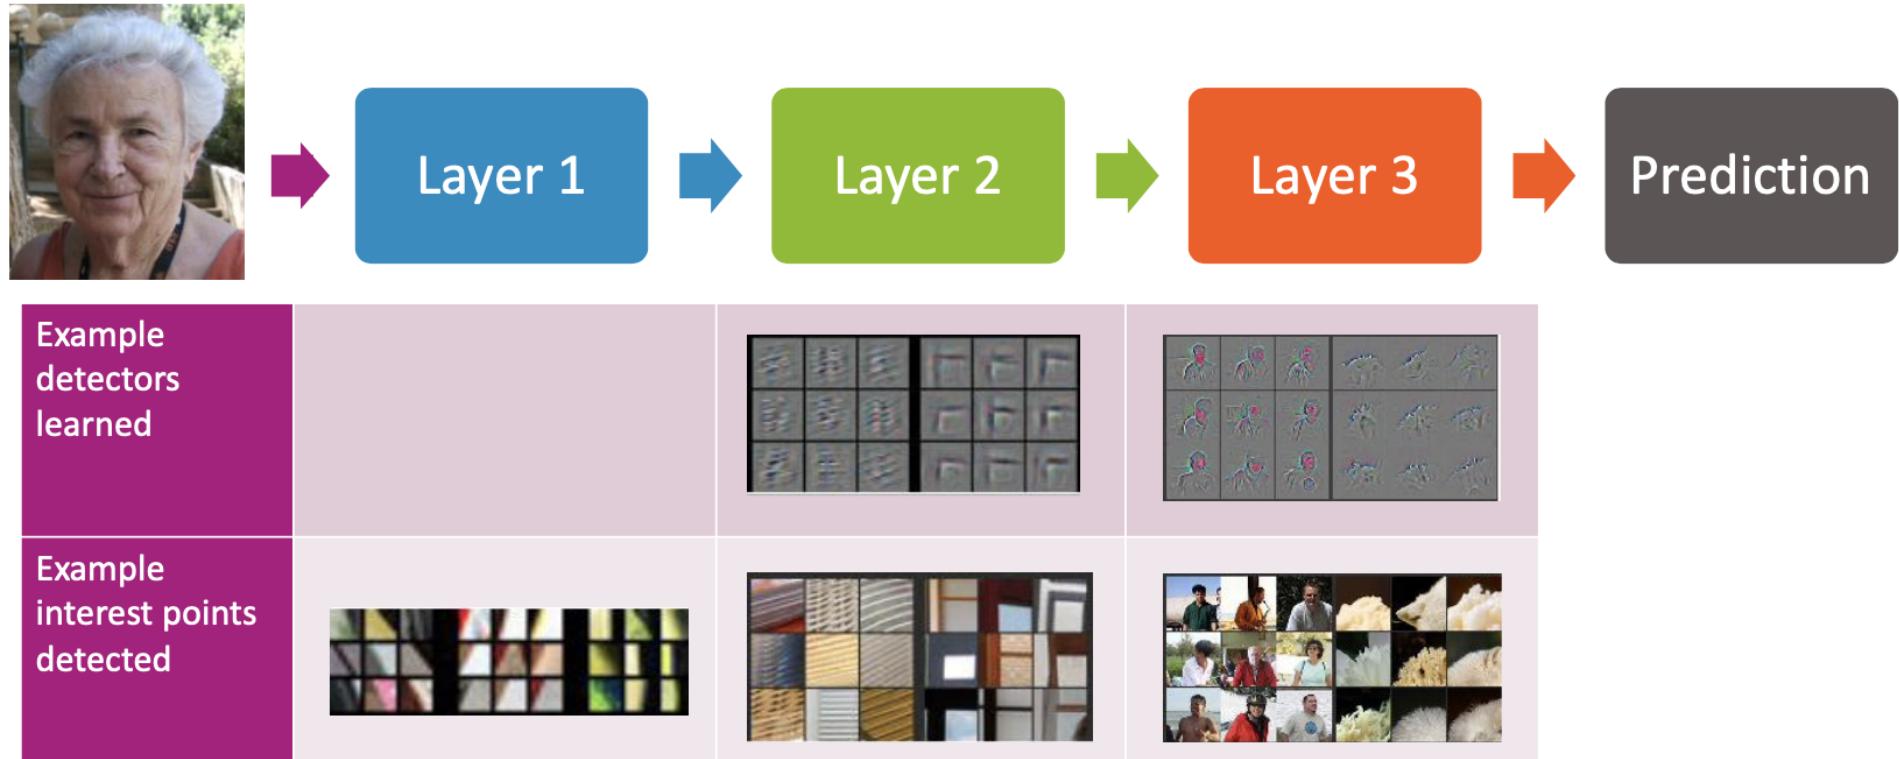


# The phenomenon of Overfitting

# Computer vision before deep learning

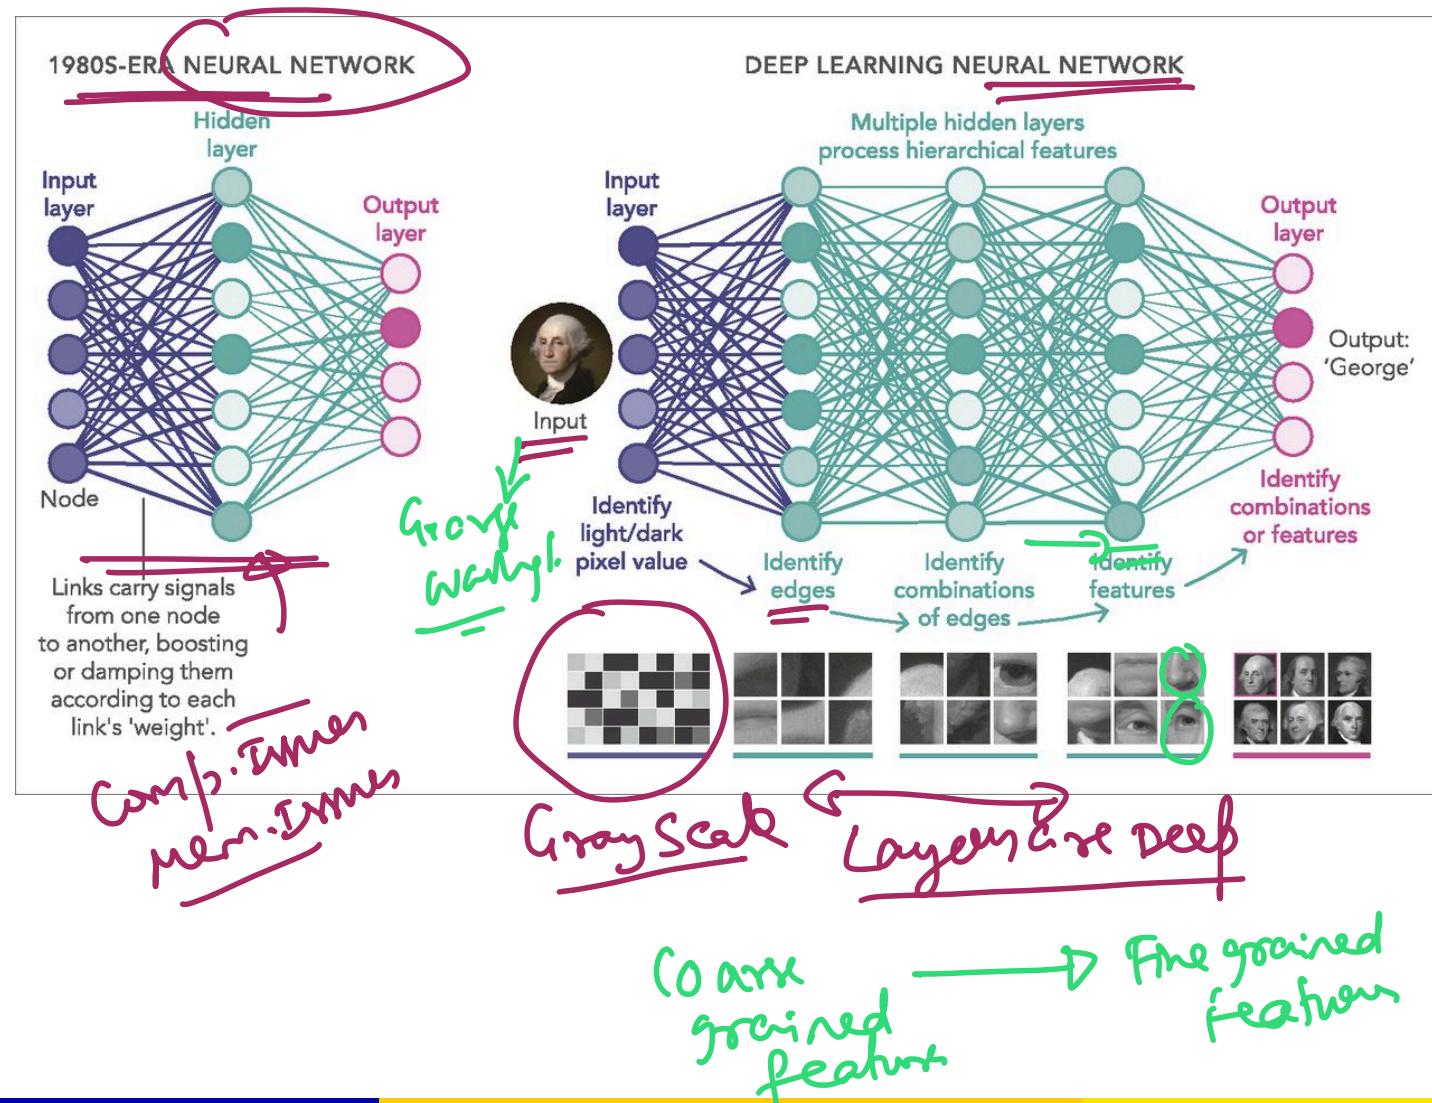


# Computer vision after deep learning

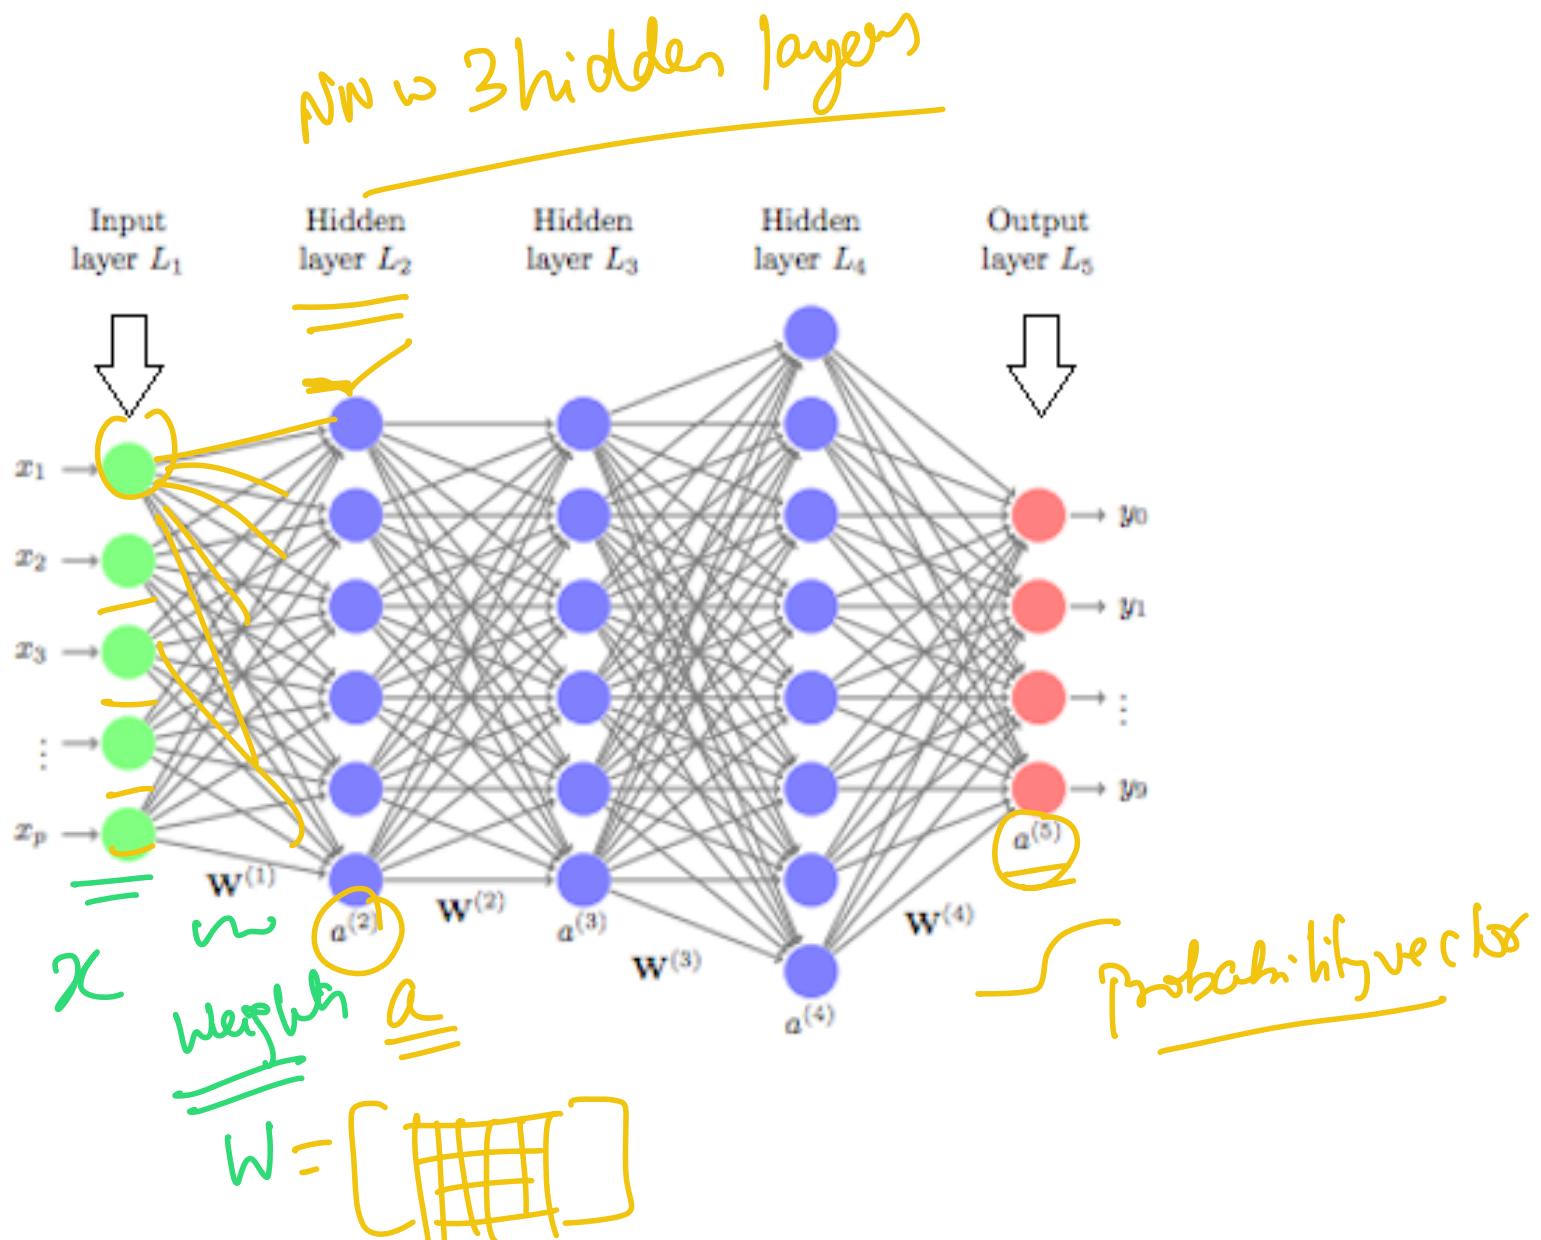


DL can extract fine-tuned features automatically

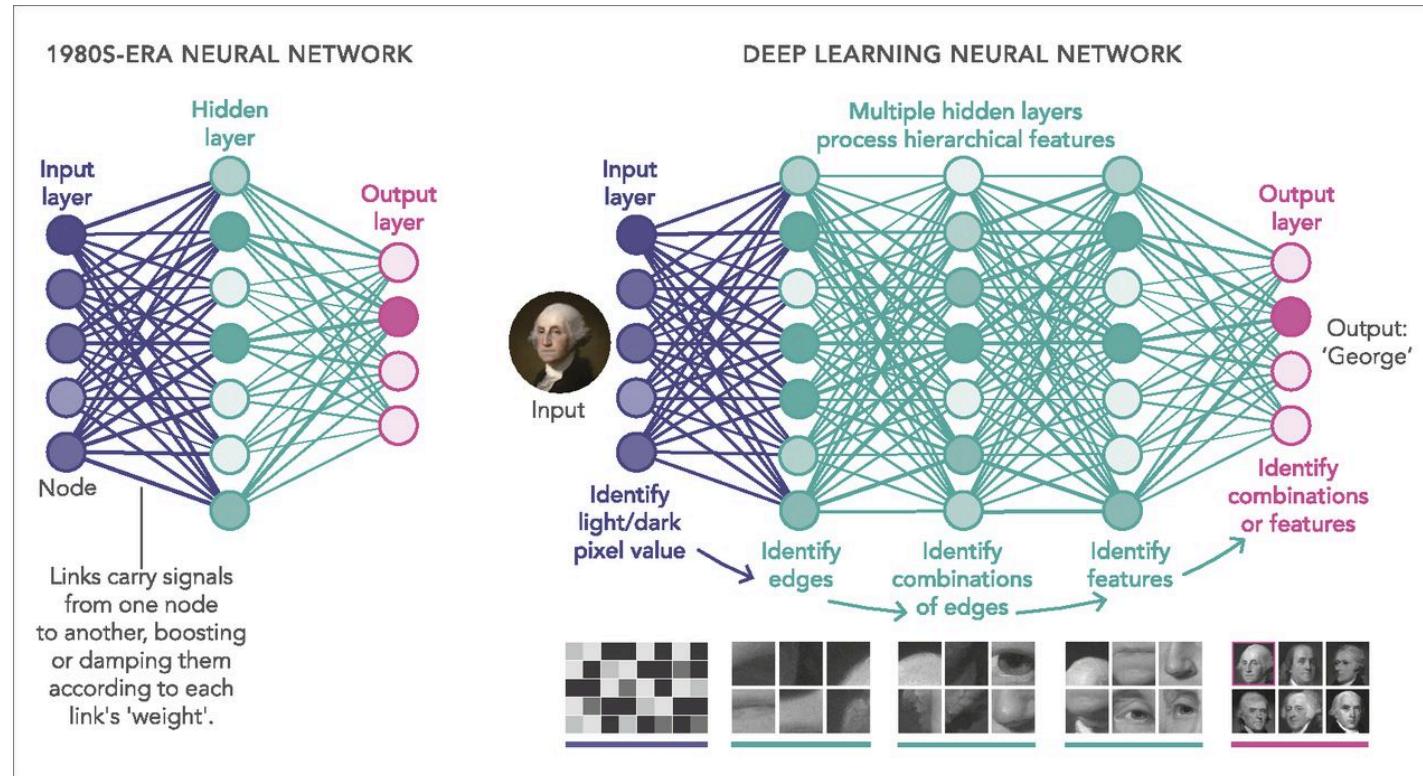
# Feed-forward Deep Learning Architecture Example



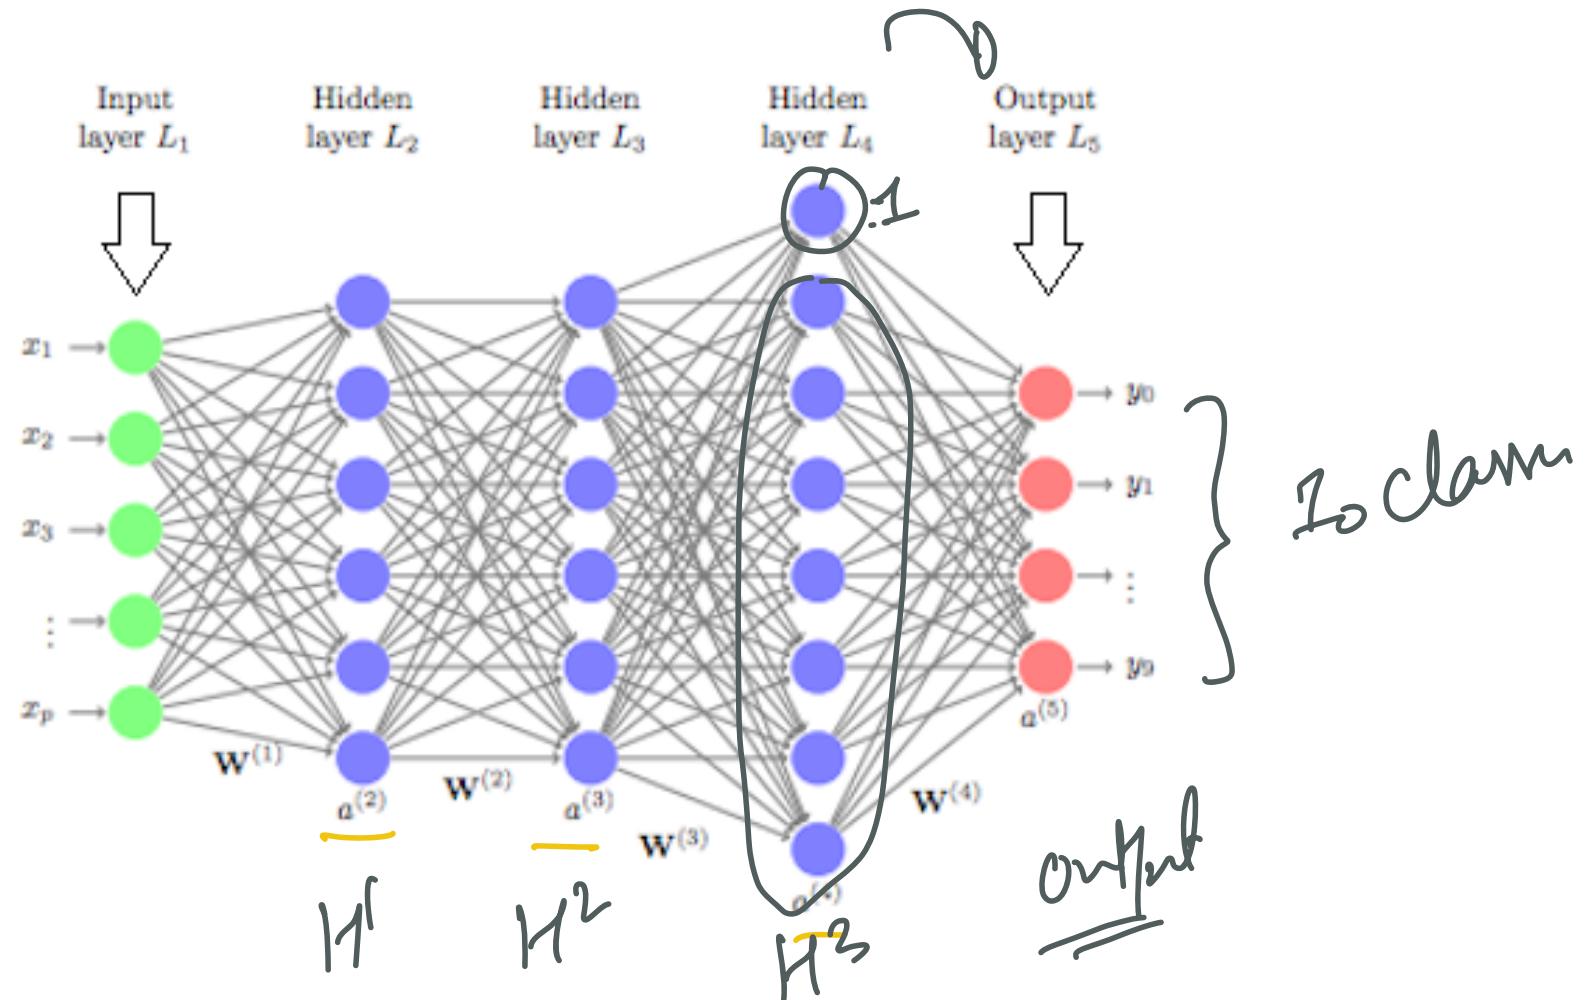
# Feed-forward Deep Learning Architecture Example



# Feed-forward Deep Learning Architecture Example



# Feed-forward Deep Learning Architecture Example

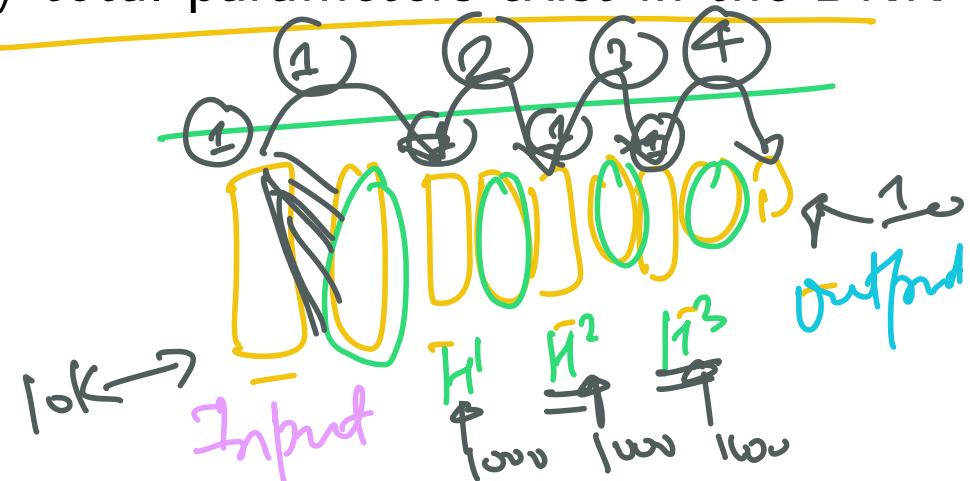


# ICE #1

Compute the number of parameters in DNN model

Consider a DNN model with 3 hidden layers where each hidden layer has 1000 neurons. Let the input layer be raw pixels from a 100x100 image and the output layer has 10 dimensions, let's say for a 10 class image classification example. How many total parameters exist in the DNN model?

- ① 10 million parameters
- ② 11 million parameters
- ③ 12 million parameters
- ④ 13 million parameters



$$\begin{aligned} \textcircled{3} &= 10K = 1M \\ \textcircled{4} &= 1000 \times 10 = 10K \end{aligned}$$

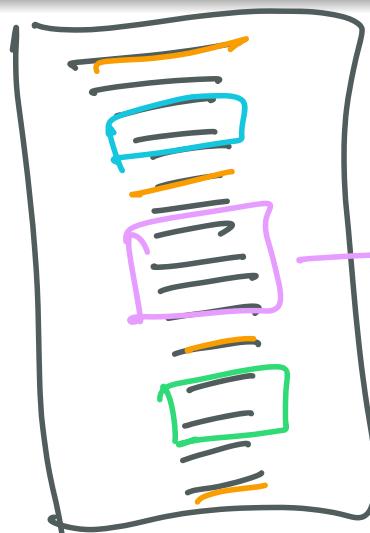
$$\begin{aligned} \textcircled{1} &= (100 \times 100) \times 1000 = 10M \\ \textcircled{2} &= 1000 \times 1000 = 1M \end{aligned}$$

$(1+2)(3+4)$   
 $= 12M + \text{change}$

# Training a DNN

SGD with mini-batch → Stochastic Gradient Descent with Mini-batch  
SGD - Main optimizer engine / Training engine for almost any ML model

SGD mini-batch is the staple diet. However there are some **learning rate schedulers** that are known to work better for DNNs - Such as Adagrad and more recently, ADAM. ADAM adapts the learning rate to each individual parameter instead of having a global learning rate.



mini-batch  
 $\approx 128/256$

→ Inform the weights  
↓  
Back-prop

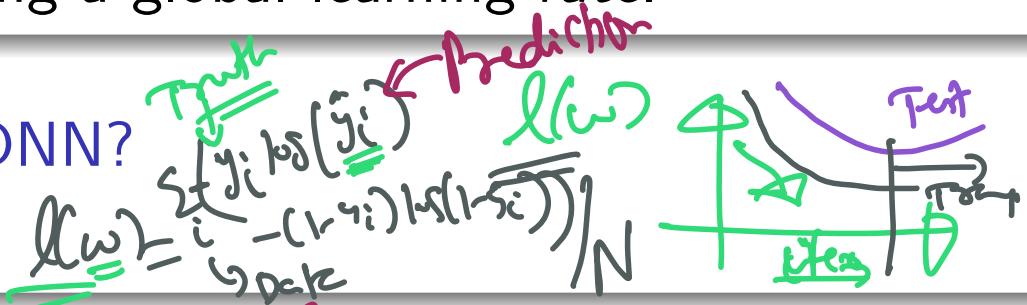
# Training a DNN

## SGD with mini-batch

SGD mini-batch is the staple diet. However there are some **learning rate schedulers** that are known to work better for DNNs - Such as Adagrad and more recently, ADAM. ADAM adapts the learning rate to each individual parameter instead of having a global learning rate.

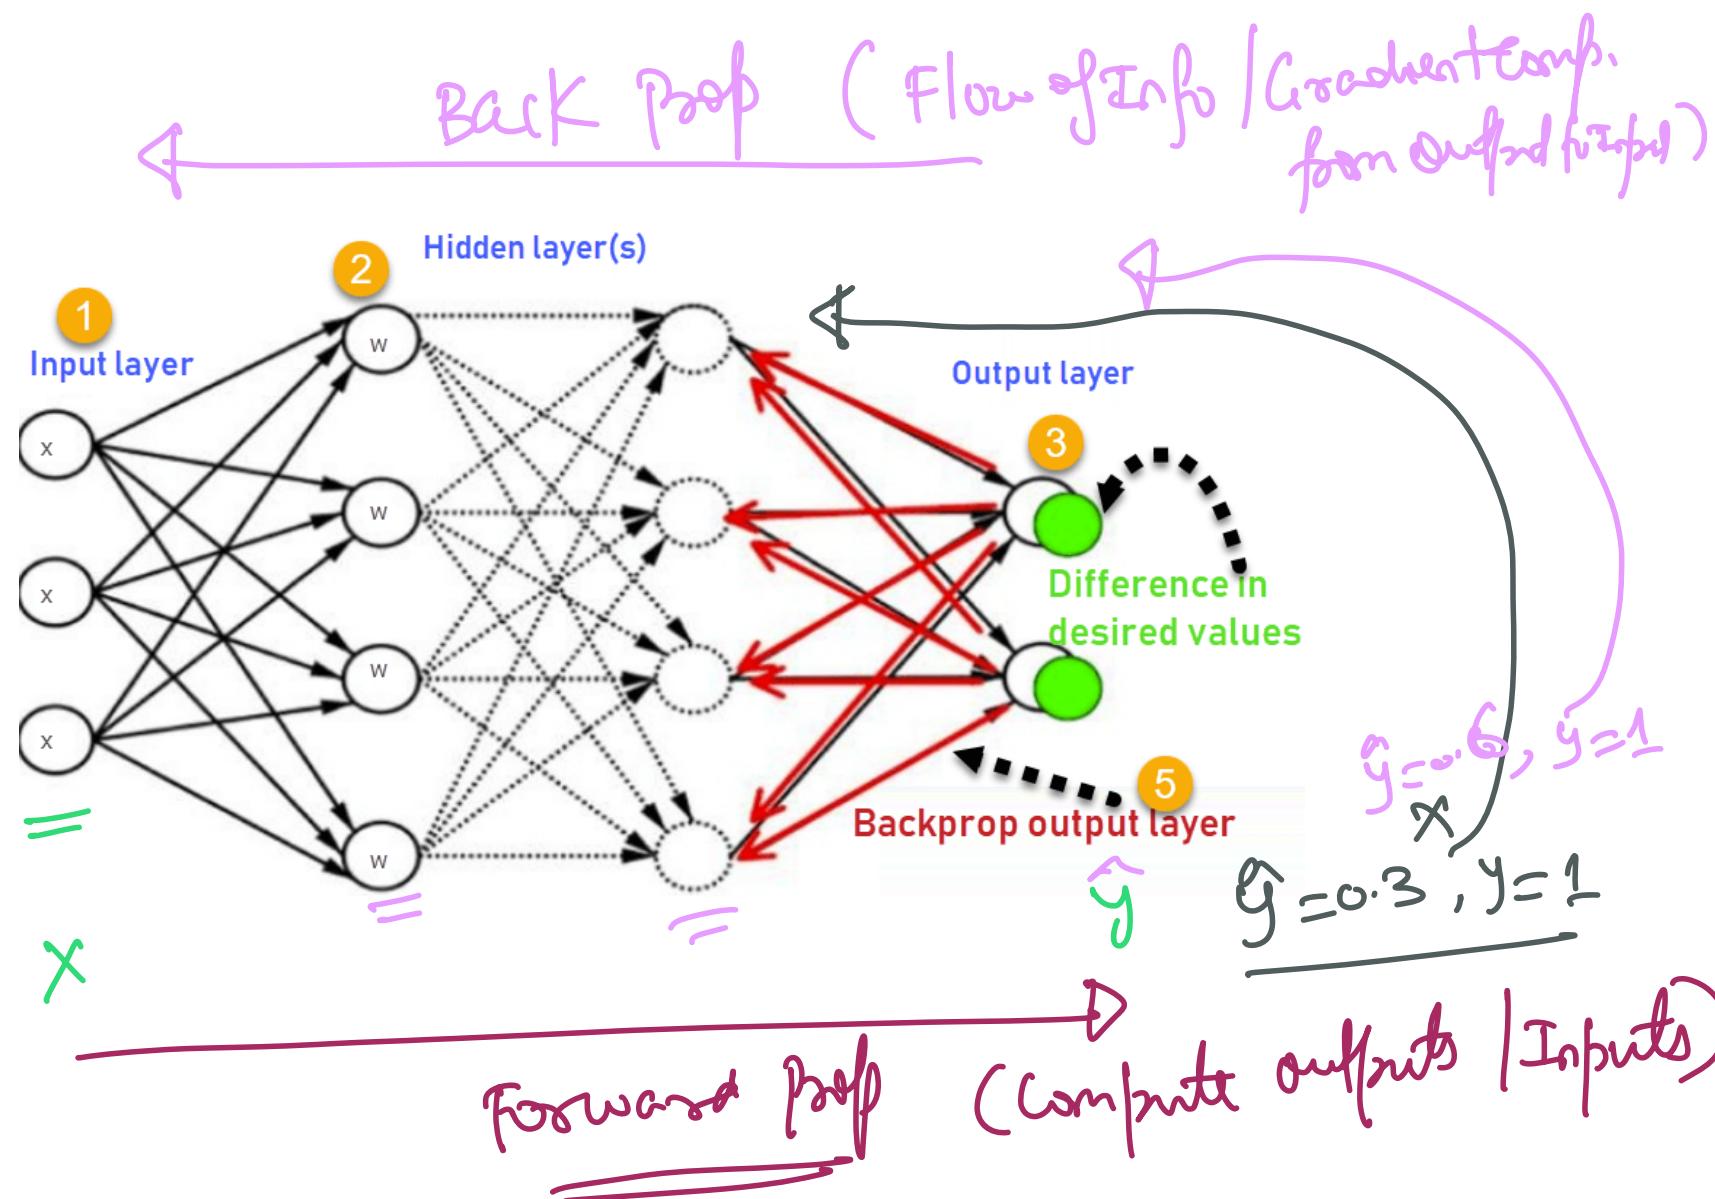
How do we compute gradient in a DNN?

Back-propagation!

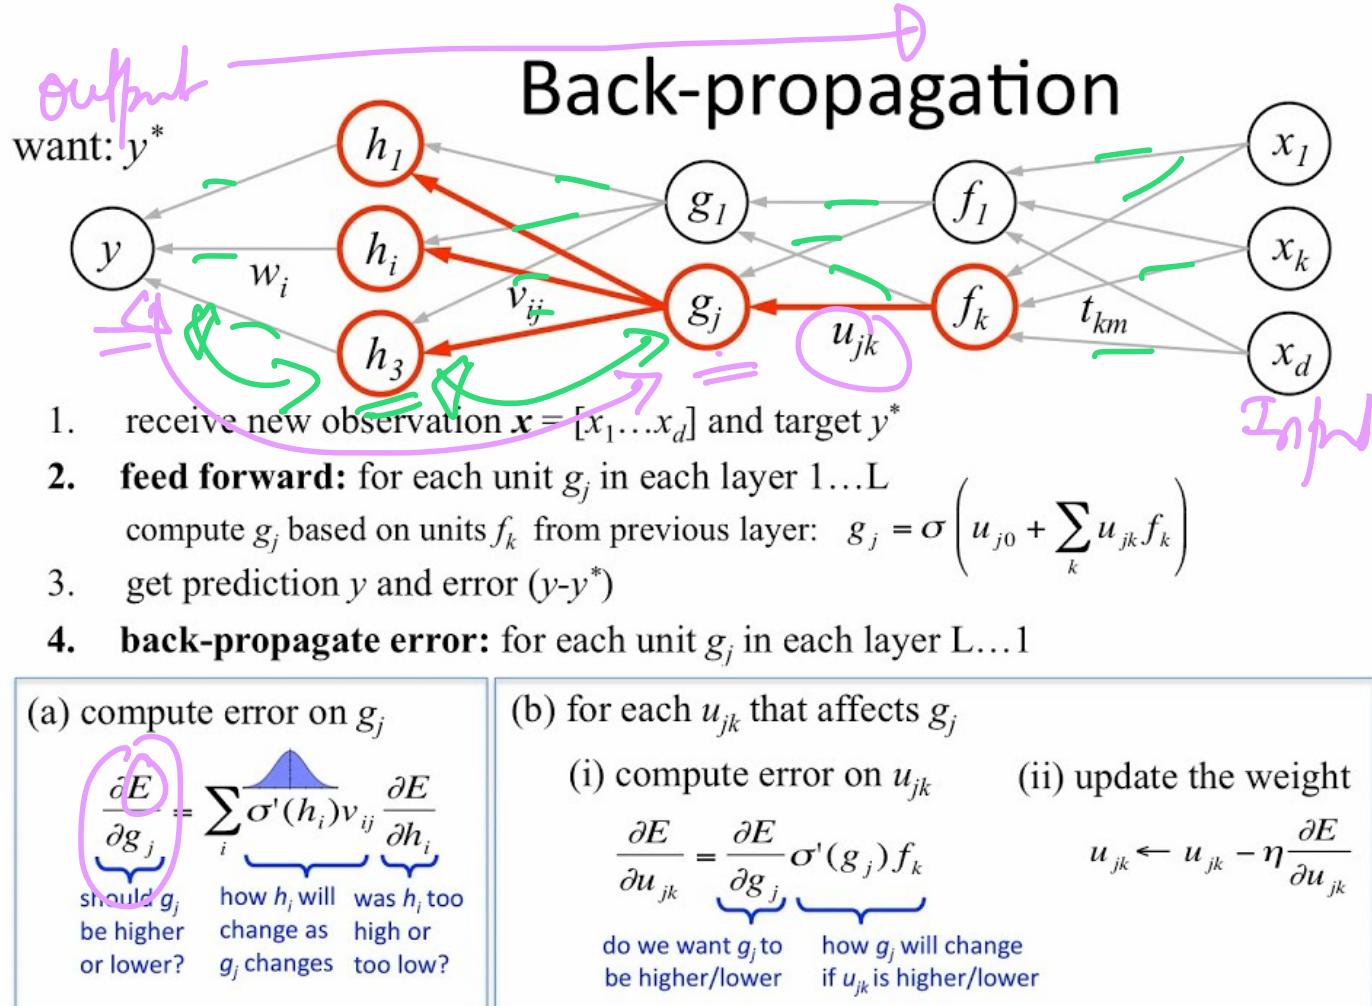


$\frac{d}{dw} l(w)$   $w^k - \text{lr. } \nabla l(w)$   $\rightarrow$  Gradient - mathematical quantity that transfers info. for learning weights in a DNN  
start  $w^0 \rightarrow w^1 \rightarrow w^2 \rightarrow \dots \rightarrow w^k$   $\rightarrow$  Trained weight

# Forward Propagation vs Back-propagation



# Back Propagation explained



Copyright © 2014 Victor Lavrenko

# Back Propagation Summary

## Back Prop

Back prop is one of the fundamental backbones of the training modules behind deep learning and beyond (including for example ChatGPT). What exactly is back prop? It is just a way to unravel gradient computation in the neural network. Back prop is how we would **compute the gradient** in a neural network.

$$f(\omega) = \omega^2$$
$$\nabla f(\omega) = 2\omega = \frac{df}{d\omega}$$
$$f(\underbrace{\omega_1, \omega_2}) = \omega_1^2 + \omega_2^2$$
$$\nabla f(\omega) = \begin{bmatrix} \frac{\partial f}{\partial \omega_1} \\ \frac{\partial f}{\partial \omega_2} \end{bmatrix} = \begin{bmatrix} 2\omega_1 \\ 2\omega_2 \end{bmatrix}$$
$$\omega = \begin{bmatrix} \omega_1 \\ \omega_2 \end{bmatrix}$$
$$= 2 \begin{bmatrix} \omega_1 \\ \omega_2 \end{bmatrix} = 2\omega$$

vector

# Back Propagation Summary

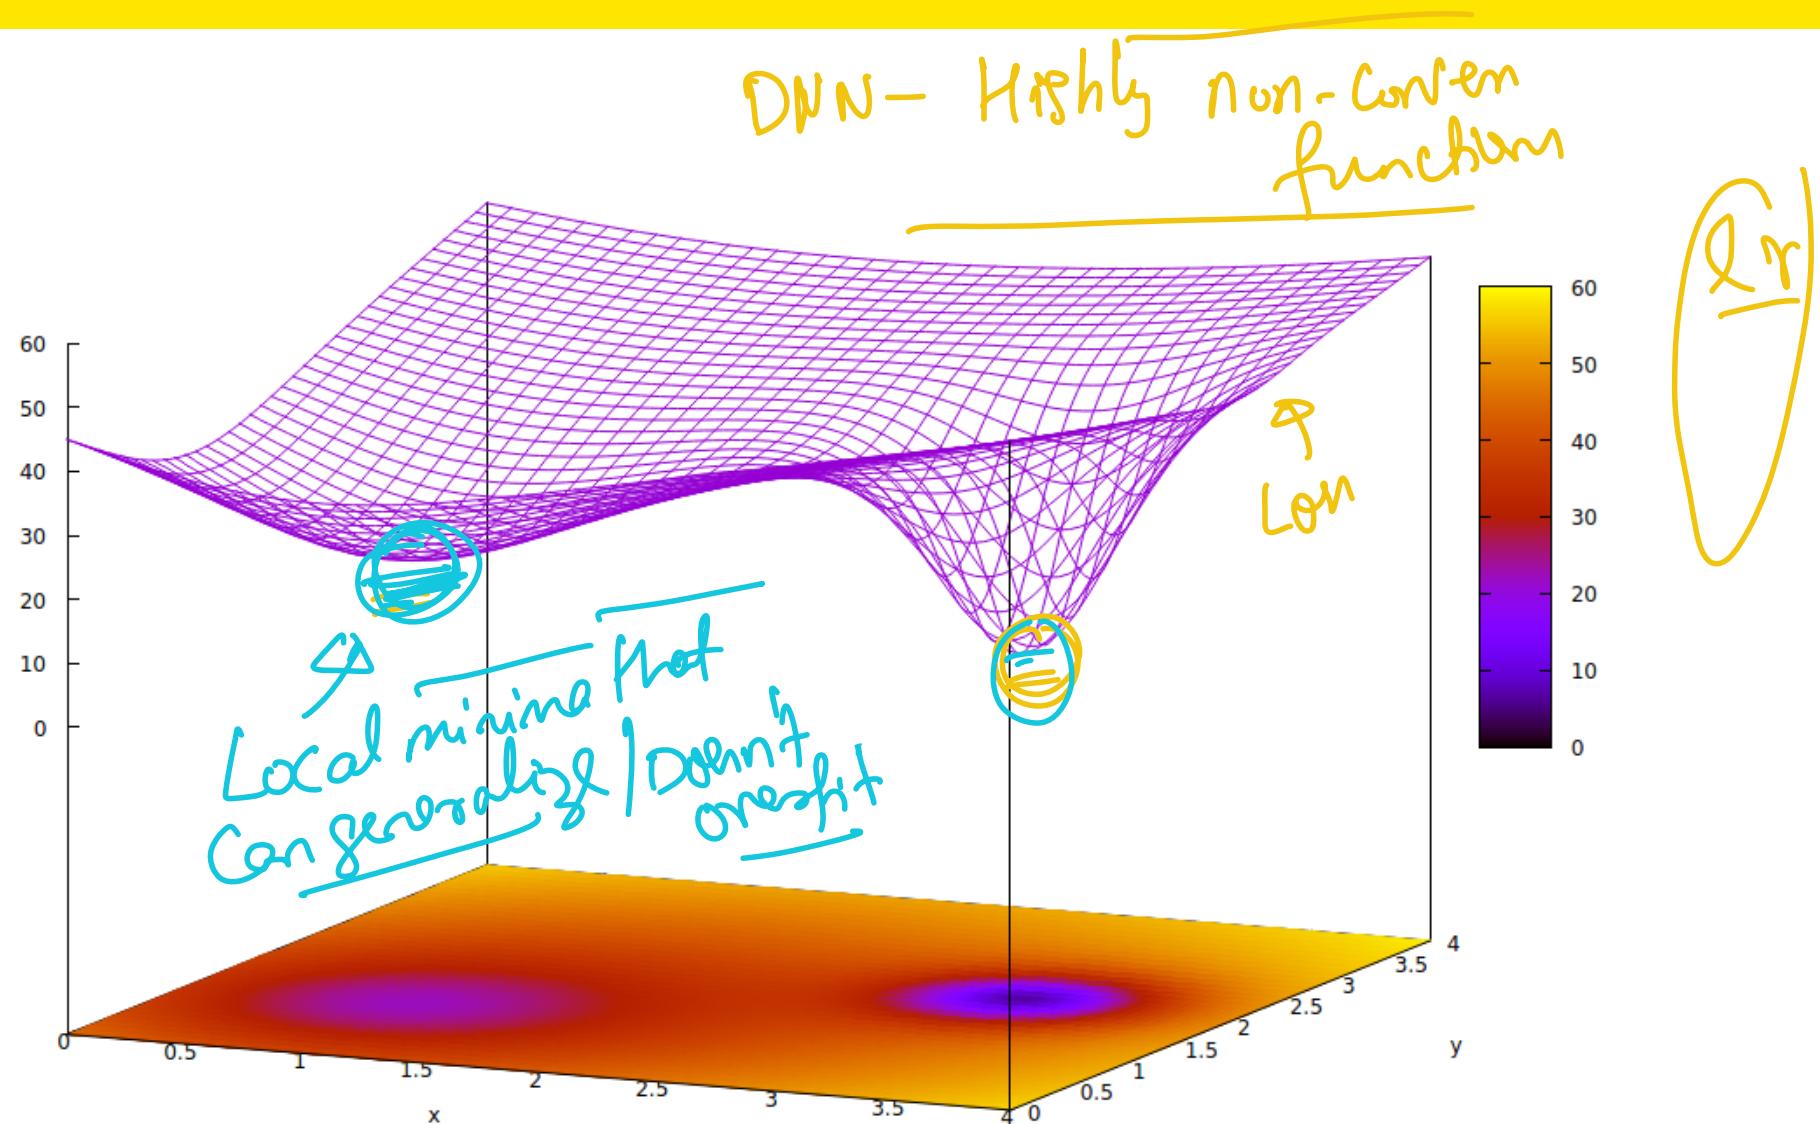
## Back Prop

Back prop is one of the fundamental backbones of the training modules behind deep learning and beyond (including for example ChatGPT). What exactly is back prop? It is just a way to unravel gradient computation in the neural network. Back prop is how we would **compute the gradient** in a neural network.

## Back Prop as information flow

It can also be thought of as flow information from the error in the output (the loss function) down to the weights. Update the weights so we don't make **this error** next time around. Back prop is a way to do **gradient descent in neural networks!**

# Good vs Bad Local minima



# Hyper-parameters in Deep Learning

ICE #2: Which of the following is not a hyper-parameter in deep learning?

- 1 Learning rate
- 2 Number of Hidden Layers
- 3 Number of neurons per hidden layer
- 4 All of the above

0.01, 0.1, 0.0001

4 Lf → For k & Training  
↓ Lf → Show & Training

# Hyper-parameters in Deep Learning

## Hyper-parameters

- ① Learning rate
- ② Number of Hidden Layers
- ③ Number of neurons per hidden layer

# Hyper-parameters in Deep Learning

## Hyper-parameters

- ① Learning rate
  - ② Number of Hidden Layers
  - ③ Number of neurons per hidden layer
  - ④ Type of non-linear activation function used
- 

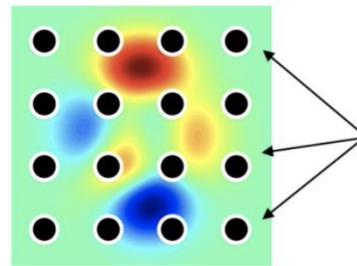
# Hyper-parameters in Deep Learning

## Hyper-parameters

- ① Learning rate
- ② Number of Hidden Layers
- ③ Number of neurons per hidden layer
- ④ Type of non-linear activation function used
- ⑤ Anything else?

# Hyper-parameter tuning methods

Grid search:



Hyperparameters  
on 2d uniform grid

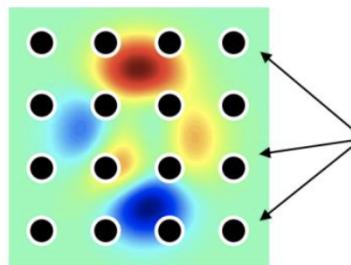
$$l_0 = \underline{0.1}$$

$$l_0 = \underline{0.01}$$

$$l_0 = \underline{\underline{0.0001}}$$

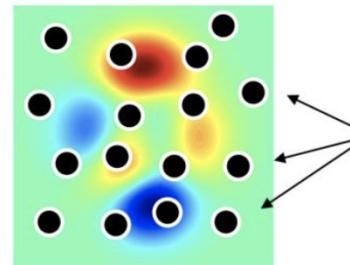
# Hyper-parameter tuning methods

Grid search:



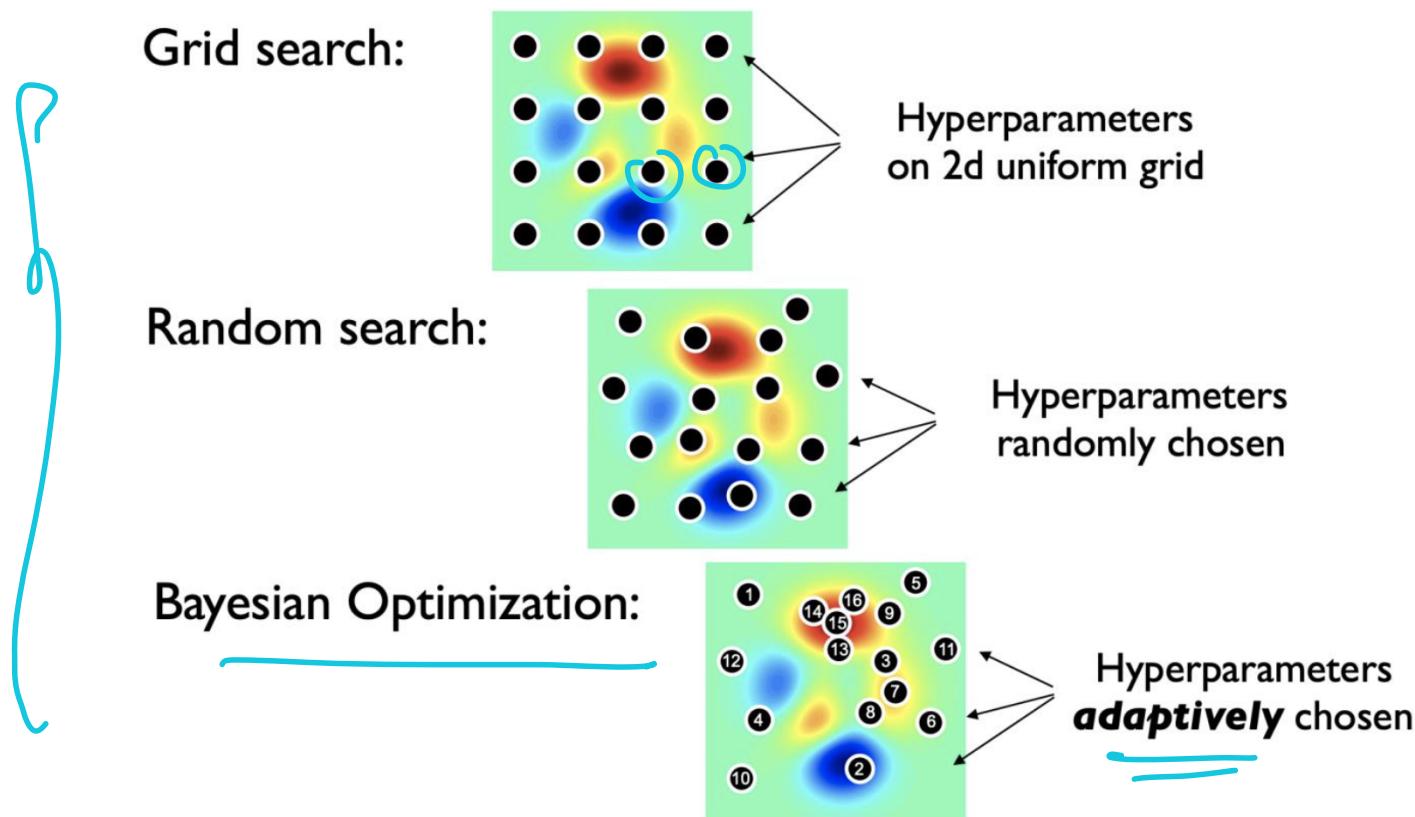
Hyperparameters  
on 2d uniform grid

Random search:



Hyperparameters  
randomly chosen

# Hyper-parameter tuning methods



# Over-fitting in DNNs

## How to handle over-fitting in DNNs

- ① A DNN model with 100 million parameters and only 100k data points or even a million data points will overfit unless we take care of over-fitting.

# Over-fitting in DNNs

## How to handle over-fitting in DNNs

- ① A DNN model with 100 million parameters and only 100k data points or even a million data points will overfit unless we take care of over-fitting.
- ② Weight regularization can help -  $\ell_1, \ell_2$

# Over-fitting in DNNs

## How to handle over-fitting in DNNs

- ① A DNN model with 100 million parameters and only 100k data points or even a million data points will overfit unless we take care of over-fitting.
- ② Weight regularization can help -  $\ell_1, \ell_2$
- ③ More common over-fitting strategy for DL?

# Over-fitting in DNNs

## How to handle over-fitting in DNNs

- ① A DNN model with 100 million parameters and only 100k data points or even a million data points will overfit unless we take care of over-fitting.
- ② Weight regularization can help -  $\ell_1, \ell_2$
- ③ More common over-fitting strategy for DL?
- ④ Dropouts!

# Over-fitting in DNNs

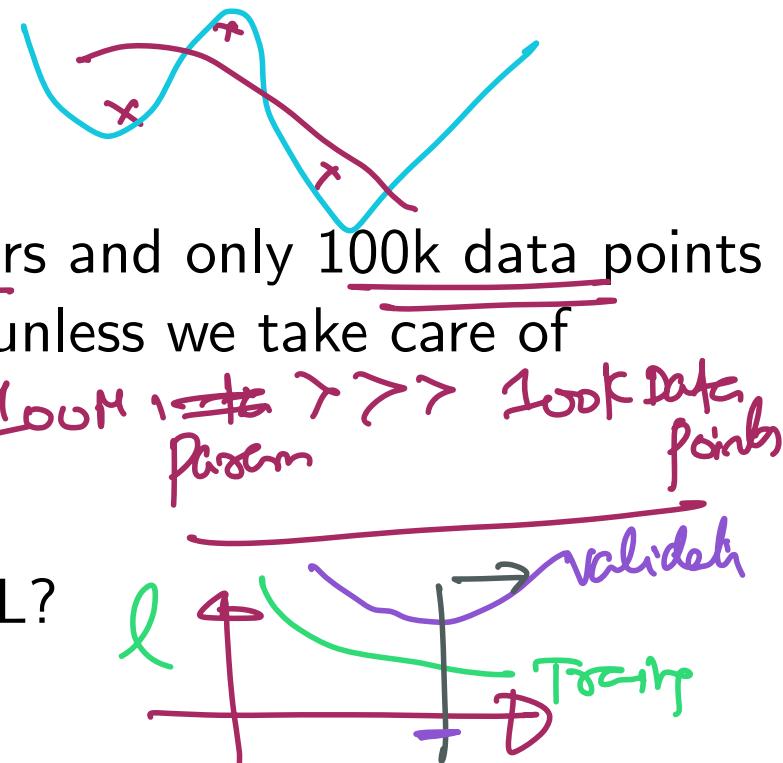
## How to handle over-fitting in DNNs

- ① A DNN model with 100 million parameters and only 100k data points or even a million data points will overfit unless we take care of over-fitting.
- ② Weight regularization can help -  $\ell_1, \ell_2$
- ③ More common over-fitting strategy for DL?
- ④ Dropouts!
- ⑤ Early stopping is also a great strategy! Stop training the DL model when the validation error starts increasing. How's this different from regular validation we were doing earlier??

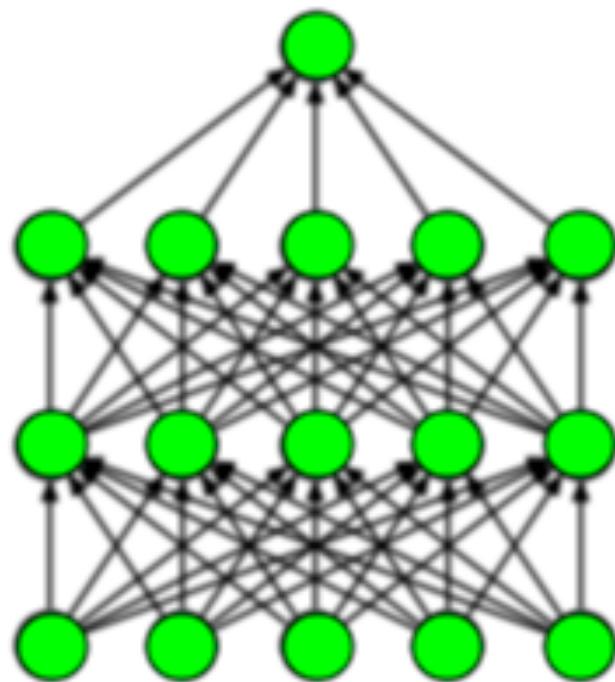
# Over-fitting in DNNs

## How to handle over-fitting in DNNs

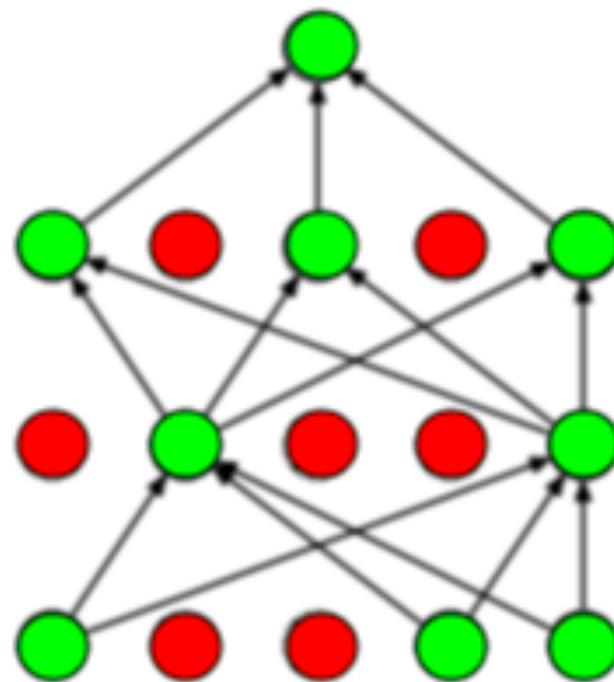
- ① A DNN model with 100 million parameters and only 100k data points or even a million data points will overfit unless we take care of over-fitting.  
*100M ~~isn't~~ >> 100k Data Points*
- ② Weight regularization can help -  $\ell_1, \ell_2$
- ③ More common over-fitting strategy for DL?
- ④ Dropouts!
- ⑤ Early stopping is also a great strategy! Stop training the DL model when the validation error starts increasing. How's this different from regular validation we were doing earlier??
- ⑥ Book by Yoshua Bengio has tons of details and great reference for Deep Learning!



# Taking care of Over-fitting: Dropouts



(a) Standard Neural Net



(b) After applying dropout.

# Tensorflow Playground Demo

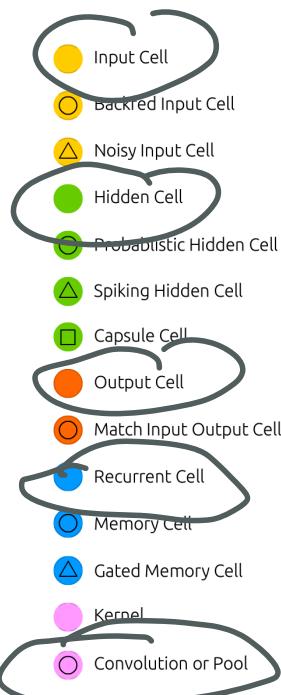
Tensorflow Playground Demo ]

# More DL Architectures

## Neural Networks Zoo

### Zoo Reference

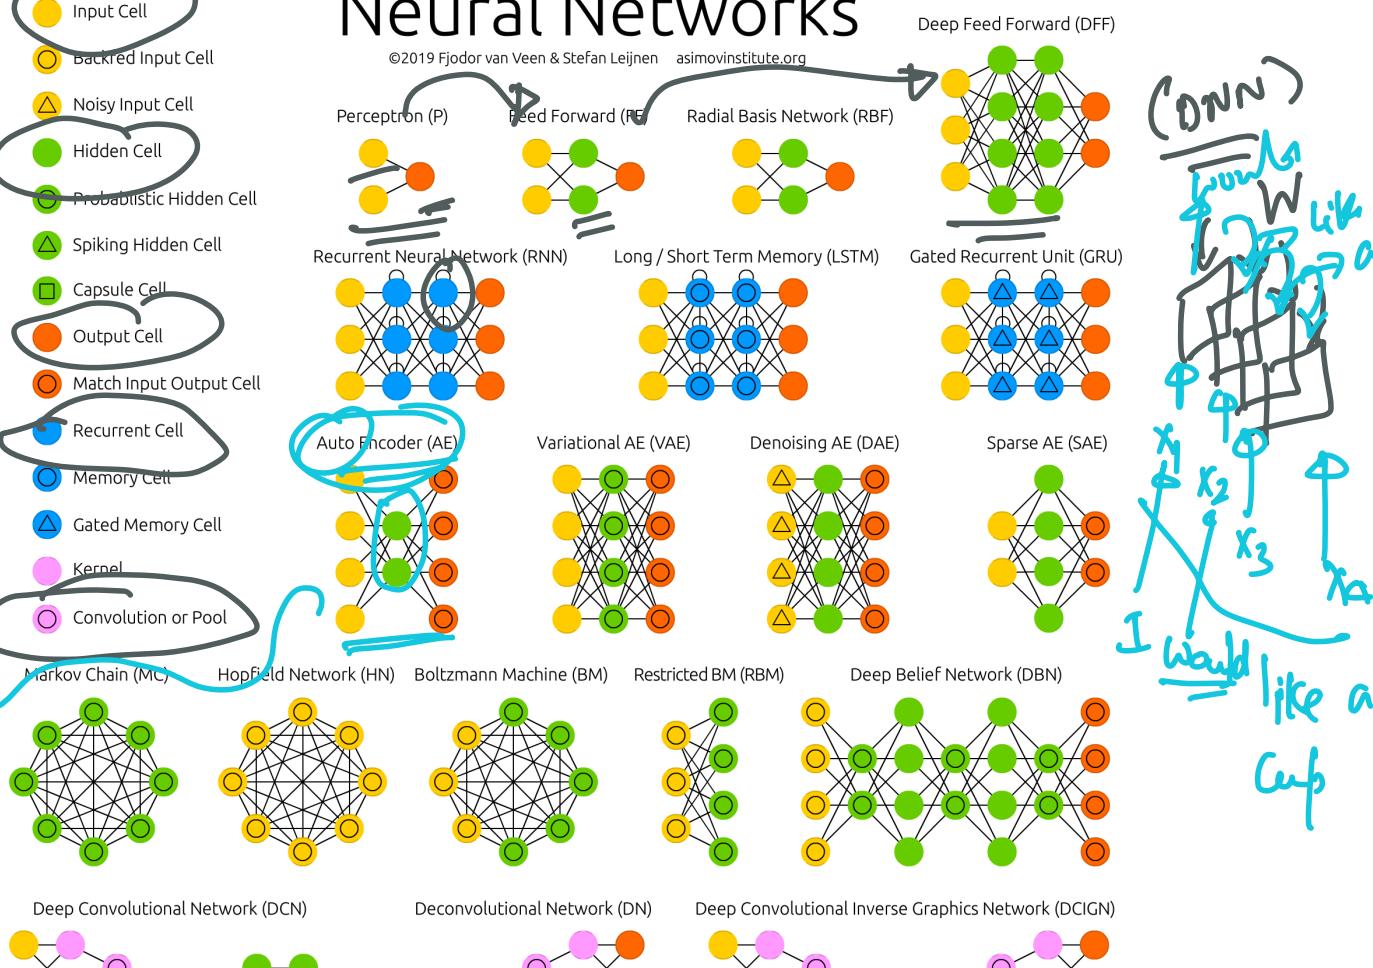
*Interpretable  
DNNs*



*Stable  
Diffusion*

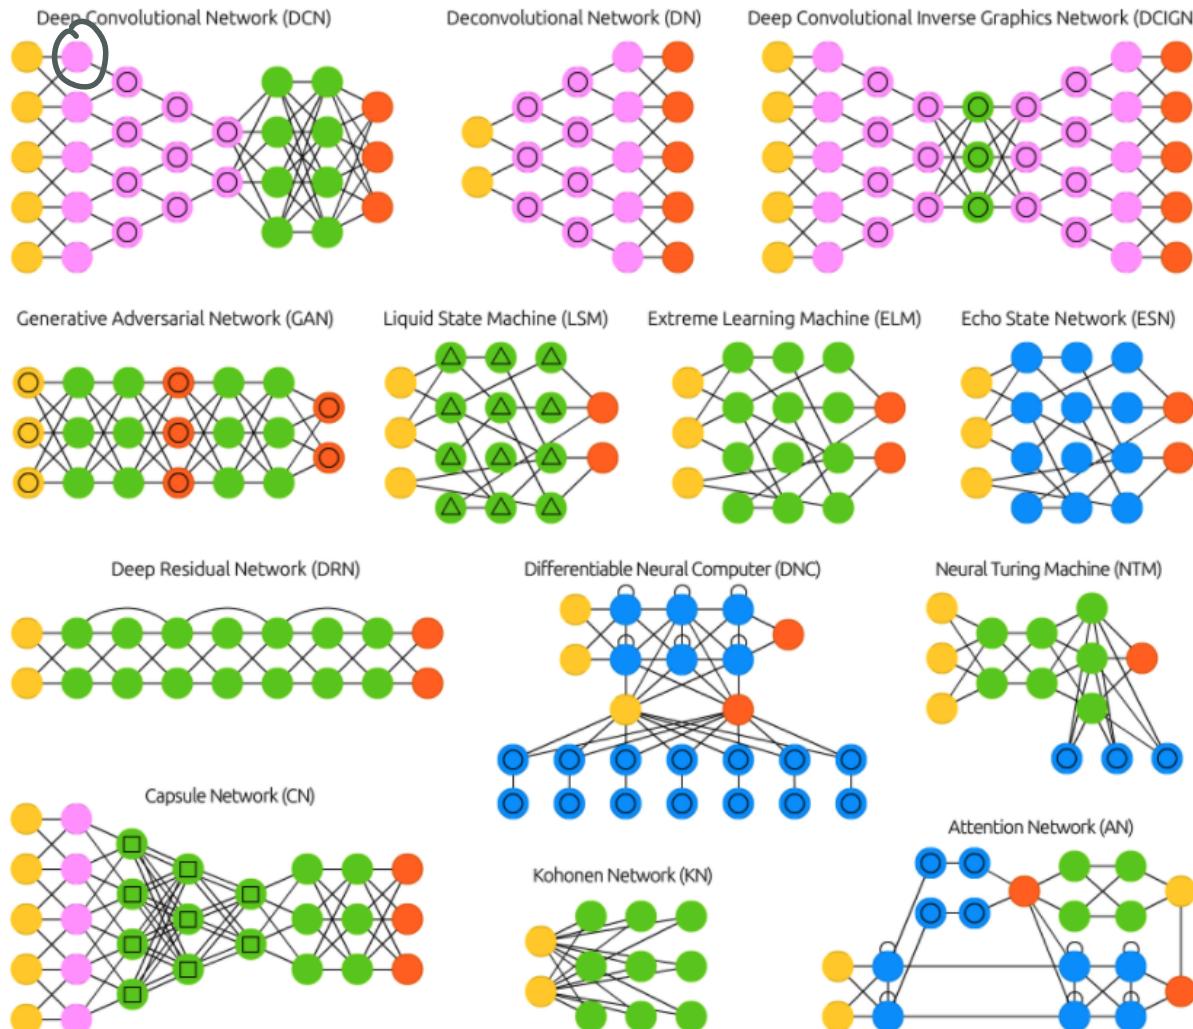
## A mostly complete chart of Neural Networks

©2019 Fjodor van Veen & Stefan Leijnen asimovinstitute.org

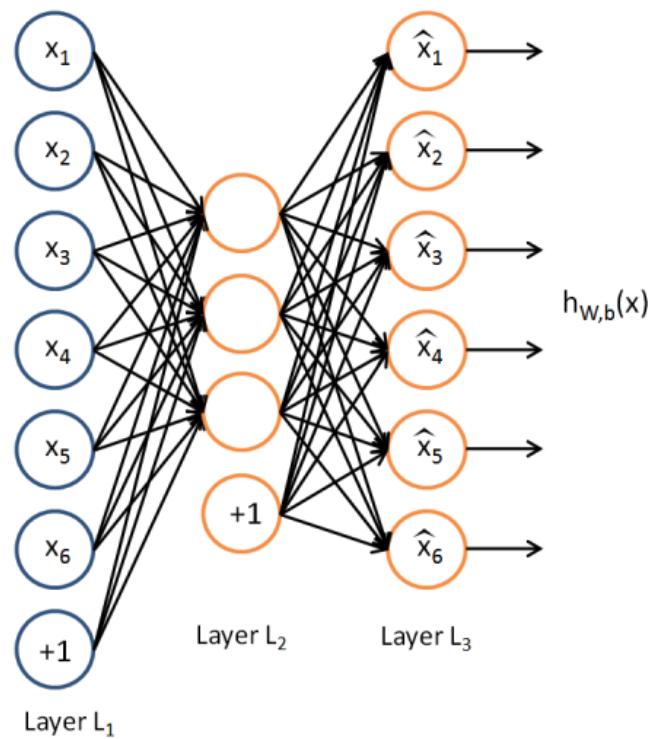


# More DL Architectures

## Neural Networks Zoo



# Auto Encoders



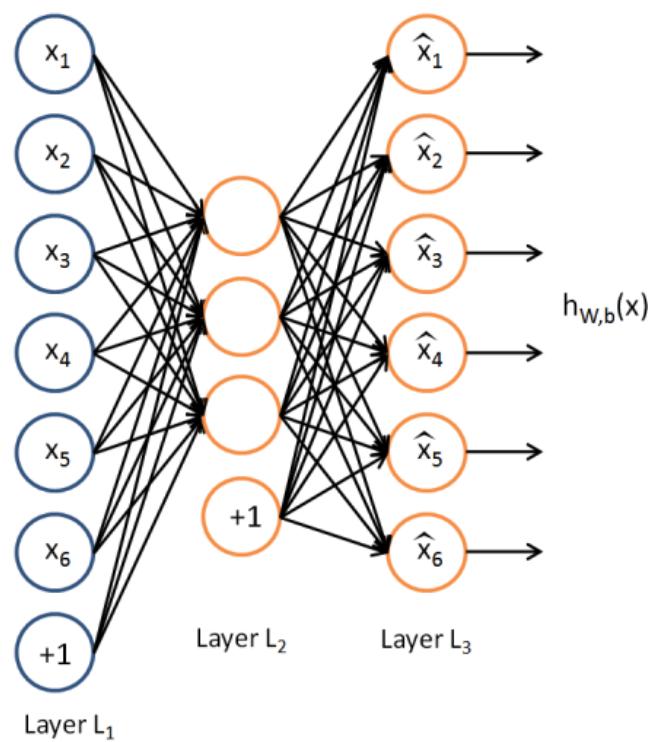
# ICE #3

## PCA vs Auto Encoder

Which of the following statements are true ?

- ① Both PCA and Auto Encoders serve the purpose of dimensionality reduction
- ② They are both linear models but one uses a neural nets architecture and the other is based on projections
- ③ PCA is robust to outliers while Auto Encoders are not
- ④ Auto Encoders are as better than Glove Embeddings to find low-dim embeddings for words

# PCA vs Auto-Encoders



# AutoEncoders and Dimensionality Reduction

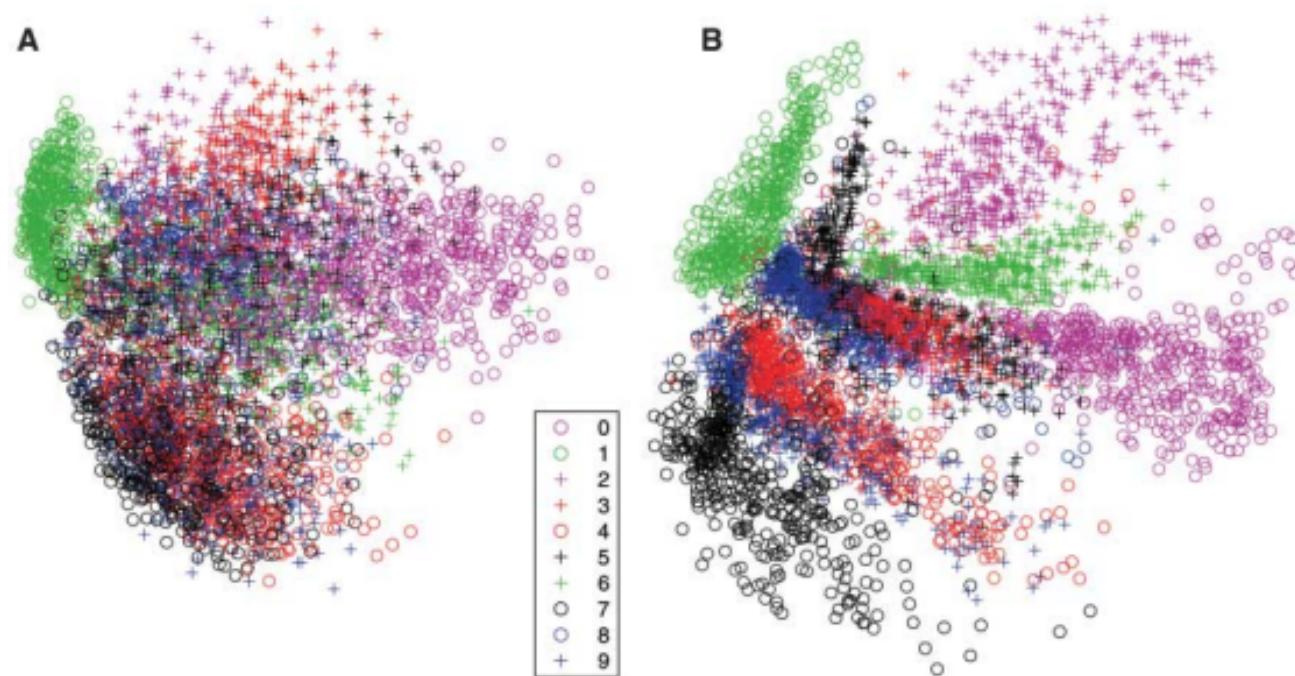
Visualization Performance

[Auto Encoder Reference Paper](#)

# AutoEncoders and Dimensionality Reduction

## Reading Reference for AE Dimensionality Reduction

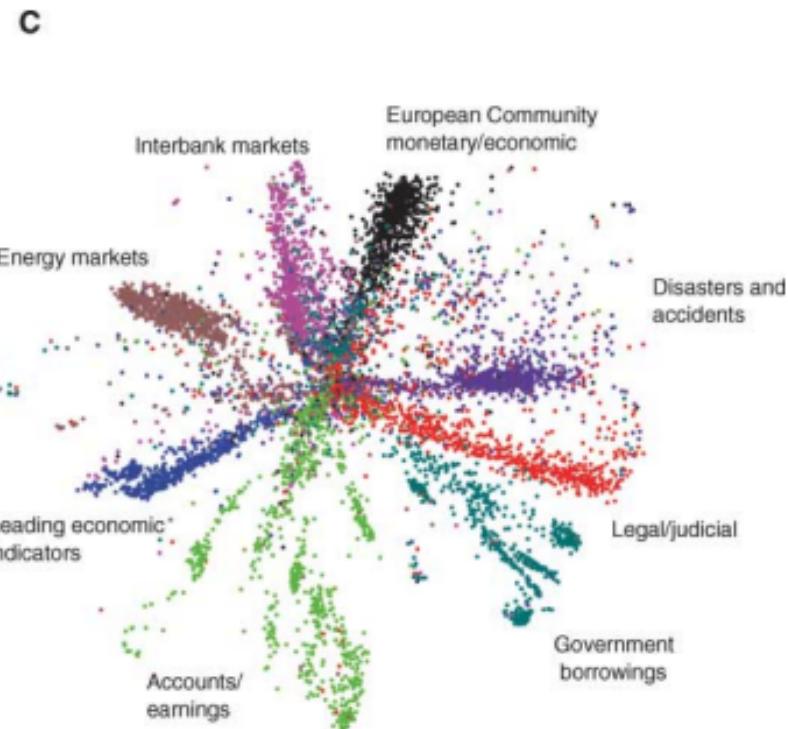
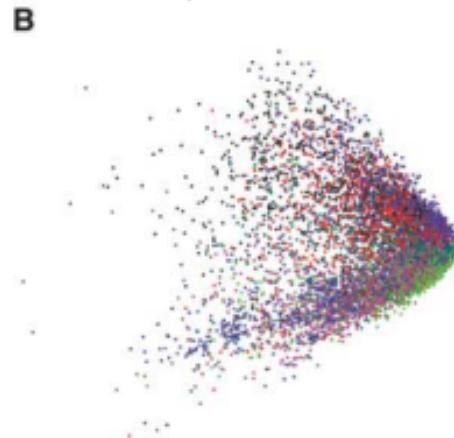
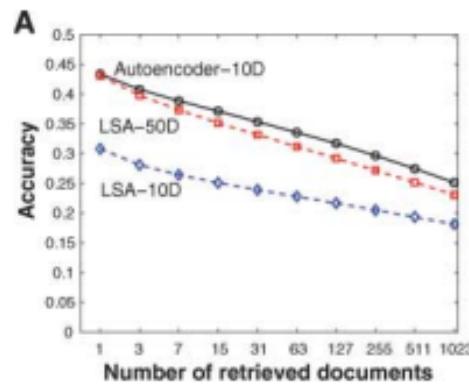
**Fig. 3.** (A) The two-dimensional codes for 500 digits of each class produced by taking the first two principal components of all 60,000 training images. (B) The two-dimensional codes found by a 784-1000-500-250-2 autoencoder. For an alternative visualization, see (8).



# AutoEncoders and Dimensionality Reduction

## Reading Reference for AE Dimensionality Reduction

**Fig. 4.** (A) The fraction of retrieved documents in the same class as the query when a query document from the test set is used to retrieve other test set documents, averaged over all 402,207 possible queries. (B) The codes produced by two-dimensional LSA. (C) The codes produced by a 2000-500-250-125-2 autoencoder.



# AutoEncoders Summary

- ① Auto-Encoders are a method for dimensionality reduction and can do better than PCA for visualization

# AutoEncoders Summary

- ① Auto-Encoders are a method for dimensionality reduction and can do better than PCA for visualization
- ② Use Neural Networks architecture and hence can encode non-linearity in the embeddings

# AutoEncoders Summary

- ① Auto-Encoders are a method for dimensionality reduction and can do better than PCA for visualization
- ② Use Neural Networks architecture and hence can encode non-linearity in the embeddings
- ③ Anything else?

# AutoEncoders Summary

- ① Auto-Encoders are a method for dimensionality reduction and can do better than PCA for visualization
- ② Use Neural Networks architecture and hence can encode non-linearity in the embeddings
- ③ Anything else?
- ④ Auto Encoders can learn convolutional layers instead of dense layers - Better for images! More flexibility!!